

capstone

September 13, 2020

Introduction

Spotify is one of the most popular music streaming service in the world. Songs of various styles and genres are published on the platform by both indie and professional artists from all over the world. As an aspirant artist who is intending to publish his/her song on the platform, one might be interested to predict how much popularity the song is going to get based on its characteristics such tempo, key or loudness. In this analysis, we're going to look into a Spotify song dataset of 160,000 songs to explore the relationship between features of a song and its popularity. A predictive model will then be built to predict how well a song would do on Spotify based on the song's characteristics.

On Spotify, beside song tracks, there are also a great amount of audiobooks, podcast or noise tracks. In this dataset, there is a good mix of different types of tracks so we're going to filter out tracks which are not songs. Only songs data entries will be used to build the predictive model.

Lastly, we're going to make a distinction between studio recorded track and live track. Our focus in this analysis is on studio recoded tracks only. Therefore, live session recording shall not be considered.

Dataset

The dataset is named 'Spotify Dataset 1921-2020, 160k+ Tracks' and was published by user Yamac Eren Ay on Kaggle. Homepage of the dataset where details can be found is accessible [here](#).

It contains more than records of more than 160,000 tracks and was collect using Spotify Web API. There are 19 columns in the dataset, description of each column is as follow:

Primary:

id: string value Unique Spotify identifier for each track

Numerical:

acousticness: float value ranges from 0 to 1

Confidence measure of whether the track is acousitc.

danceability: float value ranges from 0 to 1

Danceability describes how suitable a track is for dancing based on a combination of musical elements including tempo, rhythm stability, beat strength, and overall regularity. A value of 0.0 is least danceable and 1.0 is most danceable.

energy: float value ranges from 0 to 1

Energy is a measure from 0.0 to 1.0 and represents a perceptual measure of intensity and activity. Typically, energetic tracks feel fast, loud, and noisy.

duration_ms: integer value, ~250,000 The duration of the track in milliseconds

instrumentalness: (Ranges from 0 to 1)

Predicts whether a track contains no vocals. “Ooh” and “aah” sounds are treated as instrumental in this context. Rap or spoken word tracks are clearly “vocal”. The closer the instrumentalness value is to 1.0, the greater likelihood the track contains no vocal content. Values above 0.5 are intended to represent instrumental tracks, but confidence is higher as the value approaches 1.0.

valence: (Ranges from 0 to 1)

A measure from 0.0 to 1.0 describing the musical positiveness conveyed by a track. Tracks with high valence sound more positive (e.g. happy, cheerful, euphoric), while tracks with low valence sound more negative (e.g. sad, depressed, angry).

popularity: (Ranges from 0 to 100)

The popularity of a track is a value between 0 and 100, with 100 being the most popular. The popularity is calculated by algorithm and is based, in the most part, on the total number of plays the track has had and how recent those plays are.

tempo: (Float typically ranging from 50 to 150)

The overall estimated tempo of a track in beats per minute (BPM). In musical terminology, tempo is the speed or pace of a given piece and derives directly from the average beat duration.

liveness: (Ranges from 0 to 1)

Detects the presence of an audience in the recording. Higher liveness values represent an increased probability that the track was performed live. A value above 0.8 provides strong likelihood that the track is live.

loudness: (Float typically ranging from -60 to 0)

The overall loudness of a track in decibels (dB). Loudness values are averaged across the entire track and are useful for comparing relative loudness of tracks

speechiness: (Ranges from 0 to 1)

Speechiness detects the presence of spoken words in a track. The more exclusively speech-like the recording (e.g. talk show, audio book, poetry), the closer to 1.0 the attribute value. Values above 0.66 describe tracks that are probably made entirely of spoken words. Values between 0.33 and 0.66 describe tracks that may contain both music and speech, either in sections or layered, including such cases as rap music. Values below 0.33 most likely represent music and other non-speech-like tracks.

year: (Ranges from 1921 to 2020)

The year a track was published Dummy:

mode: (0 = Minor, 1 = Major)

explicit: (0 = No explicit content, 1 = Explicit content)

Categorical:

key: (All keys on octave encoded as values ranging from 0 to 11, starting on C as 0, C# as 1 and so on...)

artists: (List of artists mentioned)

release_date: (Date of release mostly in yyyy-mm-dd format, however precision of date may vary)

name: (Name of the track)

Let's take a look at an example record of the song "Radio Ga Ga" by "Queen". Value of 'key' field being 5 indicates that the song is in key F Major. One interesting feature of this track is 'valence' which the song has a score of 0.632. This tells us that the song is musically positive or, in another word, sounds cheerful and upbeat, which is really the case for the song "Radio Ga Ga". Another interesting metric of this song is the 'danceability' score, which it scores an impressive 0.752. This is again understandable since the song, with a simple and catchy beat, is easy to dance to.

It is important to note that these numerical metrics are given to each song using Spotify own machine learning model and aren't assigned human so we can assume that scores for different attributes of each song are given fairly consistently. In order to create an actual test case for a new track, it is important to get values for features of interest through Spotify API rather than having someone artist to assign scores to the track based on his/her feeling. More information on Spotify Web API can be found [here](#)

The key data which describe the attributes of a track such as danceability, acousticness, loudness, ... are already in numerical format and are already normalized so that will help a lot in the data wrangling process. Data field such as 'artists' or 'name', though might very well have an impact on the popularity of a song, will not be used in this model since in the scope of this analysis, we only want to focus on the musical virtues of a song.

The dataset was last updated on 19 June, 2020 so it is still highly relevant at the point of writing of this notebook. However, if this model is to be used at some point in the future, it is recommended to update dataset. Data is valid for the US region only.

```
[1]: # Import relevant libraries and read data from .csv file
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import PolynomialFeatures
from sklearn.pipeline import Pipeline
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, r2_score

spot = pd.read_csv('data.csv')
spot.head()
```

```
[1]:
```

	acousticness	artists	danceability \
0	0.995	['Carl Woitschach']	0.708
1	0.994	['Robert Schumann', 'Vladimir Horowitz']	0.379
2	0.604	['Seweryn Goszczyński']	0.749
3	0.995	['Francisco Canaro']	0.781

```

4          0.990  ['Frédéric Chopin', 'Vladimir Horowitz']          0.210

duration_ms  energy  explicit          id  instrumentalness  \
0      158648  0.1950          0  6KbQ3uYMLKb5jDxLF7wYDD          0.563
1      282133  0.0135          0  6KuQTIu1KoTTkLXKrwLLPV          0.901
2      104300  0.2200          0  6L63VWOPibdM1HDSBoqnoM          0.000
3      180760  0.1300          0  6M94FkXd15sOAOQYRnWPN8          0.887
4      687733  0.2040          0  6N6tiFZ9vLTSOIxkj8qKrd          0.908

key  liveness  loudness  mode  \
0   10    0.1510  -12.428    1
1    8    0.0763  -28.454    1
2    5    0.1190  -19.924    0
3    1    0.1110  -14.734    0
4   11    0.0980  -16.829    1

name  popularity  release_date  \
0      Singende Bataillone 1. Teil          0          1928
1  Fantasiestücke, Op. 111: Più tosto lento          0          1928
2      Chapter 1.18 - Zamek kaniowski          0          1928
3  Bebamos Juntos - Instrumental (Remasterizado)          0  1928-09-25
4  Polonaise-Fantaisie in A-Flat Major, Op. 61          1          1928

speechiness  tempo  valence  year
0      0.0506  118.469  0.7790  1928
1      0.0462   83.972  0.0767  1928
2      0.9290  107.177  0.8800  1928
3      0.0926  108.003  0.7200  1928
4      0.0424   62.149  0.0693  1928

```

```

[2]: # For the example in 'Dataset' section
spot[spot['name']=='Radio Ga Ga']

```

```

[2]:    acoustictness  artists  danceability  duration_ms  energy  explicit  \
84438          0.151  ['Queen']          0.752        348173    0.375          0

          id  instrumentalness  key  liveness  loudness  \
84438  2jAc9KIQ9XoZxkydXh3MVh          0.000481    5    0.143  -12.966

mode          name  popularity  release_date  speechiness  tempo  \
84438    1  Radio Ga Ga          52  1984-02-27          0.0358  112.415

valence  year
84438    0.632  1984

```

Methodology

First we're going to do some exploratory analysis to have a preliminary understanding of the dataset.

Next, we're going to clean up the dataset by filtering out data entries which are not valid, followed by data normalization. Lastly, we're going to select which attributes to use as independent variables in building the predictive model. Since the variable we're trying to predict is a continuous value, we're going to use a Regression model, more details below.

```
[3]: print(spot.shape)
      print(spot.dtypes)
```

```
(169909, 19)
acousticness      float64
artists           object
danceability       float64
duration_ms       int64
energy            float64
explicit          int64
id                object
instrumentalness   float64
key               int64
liveness          float64
loudness          float64
mode              int64
name              object
popularity         int64
release_date      object
speechiness        float64
tempo             float64
valence           float64
year              int64
dtype: object
```

```
[4]: spot.describe()
```

```
[4]:
```

	acousticness	danceability	duration_ms	energy \
count	169909.000000	169909.000000	1.699090e+05	169909.000000
mean	0.493214	0.538150	2.314062e+05	0.488593
std	0.376627	0.175346	1.213219e+05	0.267390
min	0.000000	0.000000	5.108000e+03	0.000000
25%	0.094500	0.417000	1.710400e+05	0.263000
50%	0.492000	0.548000	2.086000e+05	0.481000
75%	0.888000	0.667000	2.629600e+05	0.710000
max	0.996000	0.988000	5.403500e+06	1.000000

	explicit	instrumentalness	key	liveness \
count	169909.000000	169909.000000	169909.000000	169909.000000
mean	0.084863	0.161937	5.200519	0.206690
std	0.278679	0.309329	3.515257	0.176796
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	2.000000	0.098400


50%	0.000000	0.000204	5.000000	0.135000
75%	0.000000	0.086800	8.000000	0.263000
max	1.000000	1.000000	11.000000	1.000000

	loudness	mode	popularity	speechiness \
count	169909.000000	169909.000000	169909.000000	169909.000000
mean	-11.370289	0.708556	31.556610	0.094058
std	5.666765	0.454429	21.582614	0.149937
min	-60.000000	0.000000	0.000000	0.000000
25%	-14.470000	0.000000	12.000000	0.034900
50%	-10.474000	1.000000	33.000000	0.045000
75%	-7.118000	1.000000	48.000000	0.075400
max	3.855000	1.000000	100.000000	0.969000

	tempo	valence	year
count	169909.000000	169909.000000	169909.000000
mean	116.948017	0.532095	1977.223231
std	30.726937	0.262408	25.593168
min	0.000000	0.000000	1921.000000
25%	93.516000	0.322000	1957.000000
50%	114.778000	0.544000	1978.000000
75%	135.712000	0.749000	1999.000000
max	244.091000	1.000000	2020.000000

We can see there is no missing data across the fields

```
[5]: bins = np.arange(0, 100, 1)
spot.plot(kind='hist', y='popularity', bins=bins)
plt.show()
print(spot.popularity.value_counts())
```




```
0      27357
42      3280
43      3120
40      3061
44      3054
```

```
...
93         3
97         1
96         1
99         1
100        1
```

Name: popularity, Length: 100, dtype: int64

there are 27357 tracks with 0 popularity score, which is a disproportionate amount as show in the distribution graph below Let's investigate the reasons for this 0 rating, keepi in mind that popularity is calculated mostly using counts of listen and how recent those listens are

```
[6]: spot_zero = spot[spot.popularity == 0]
     bins = np.arange(1920, 2021, 1)
     spot_zero.plot(kind='hist', y='year', bins=bins)
     plt.show()
```



It's clear that the majority of the 0-popularity tracks are tracks published before 1960 these songs are probably too old for the current user base (largely millenials) to know about Therefore, these songs don't receive enough plays to get a popularity score. Some of these songs might become more popular if somehow the user base decide to give all these old tracks a try, however there's no way to know for sure. Due to this uncertainty, it's best to drop these records because there is no relationship between popularity score and a track's attributes

```
[7]: print(spot[spot.popularity==0][spot.year==2020][['artists']].squeeze().unique())
```

```
["['Tame Impala']" "['Morat']" "['Morat', 'Juanes']" "['Summer Walker']"
"['Usher', 'Tyga']" "['Wisin & Yandel', 'Anthony Santos']"
"['Wisin & Yandel', 'Chris Brown', 'T-Pain']" "['Wisin & Yandel']"
"['Wisin & Yandel', 'Jennifer Lopez']" "['Arijit Singh']"
"['Arijit Singh', 'Shadab Faridi']" "['Arijit Singh', 'Shreya Ghoshal']"
"['Flo Rida']" "['Jack Johnson']" "['Ne-Yo']" "['Becky Hill']"
"['Alejandro Fernández', 'Christina Aguilera']"
"['Alejandro Fernández', 'Morat']" "['Luis Fonsi']"
"['Luis Fonsi', 'Aleks Syntek', 'Noel Schajris', 'David Bisbal']"
"['Luis Fonsi', 'Demi Lovato']" "['Stevie Wonder']" "['Dean Lewis']"
"['Johann Sebastian Bach', 'Lucas Jussen', 'Arthur Jussen']"
"['Marc Anthony']" "['Juanes']" "['Juanes', 'Nelly Furtado']"
"['Alejandro Sanz']" "['Jax Jones', 'Ina Wroldsen']"
"['Cali Y El Dandee', 'Sebastian Yatra']" "['Bryan Adams']"
"['Dire Straits']"
"['J Balvin', 'Yandel', 'Farruko', 'Nicky Jam', 'De La Ghetto', 'Daddy Yankee',
'Zion', 'Arcangel']"
```



```
"['J Balvin']" "['J Balvin', 'Pharrell Williams', 'BIA', 'Sky']"
"['J Balvin', 'Jowell & Randy']" "['Lady Gaga']" "['Future']"
"['A Boogie Wit da Hoodie']" "['Maroon 5']"
"['KAROL G', 'Bad Bunny', 'Quavo']" "['KAROL G', 'Ozuna']"
"['Sebastian Yatra']" "['Sebastian Yatra', 'Wisin', 'Nacho']"
"['Sebastian Yatra', 'Cosculluela', 'Cali Y El Dandee']"
"['Black Eyed Peas']" "['G-Eazy']" "['6ix9ine', 'Nicki Minaj']"
"['Swae Lee']" "['Don Omar']" "['ROSALÍA']" "['Ellie Goulding']"]
```

<ipython-input-7-76ff31b8f52d>:1: UserWarning: Boolean Series key will be
reindexed to match DataFrame index.

```
print(spot[spot.popularity==0][spot.year==2020][['artists']].squeeze().unique())
```

There's a small group of tracks published recently but also have a 0 popularity score A quick query shows that these tracks are by relatively popular artists so the reason for 0 popularity score might be because these tracks were released too recently that they haven't got enough time to gather enough plays to get a score Same as above, we're going to drop these records.

```
[8]: spot[spot.popularity == 0].index
spot.drop(index=spot[spot.popularity == 0].index, inplace=True)
# drop rows where popularity score = 0
```

According to Spotify API documentation, 'liveness' is defined as detection level of the presence of audience Higher 'liveness' values indicate a higher probability that a certain track is a record of a live session. A 'liveness' score of above 0.8 presents a strong likelihood that a track is live. Since our model is meant to provide prediction for a studio recorded track, we're going to filter and discard tracks which are highly likely to be live records.

```
[9]: print((spot.liveness > 0.8).value_counts())
# There are about 3000 tracks in this dataset which are highly likely to be
↳ live records
spot.drop(index=spot[spot.liveness > 0.8].index, inplace=True)
```

```
False    139366
True       3186
Name: liveness, dtype: int64
```

We're also not going to use 'liveness' as one of the predicting variables since it's not an attribute which an artist consciously control when creating a track. As far as we're concerned, 'liveness' helps to classify whether a track is live or not and we're not concerned with how whether a track is live or not affects its popularity on Spotify.

'speechiness' above 0.66 is considered non-musical so therefore we're going to filter out those tracks as well unlike 'liveness', 'speechiness' does characterize a track rather than simply classify whether a track is a podcast/talkshow or a song For example, with 'speechiness' somewhere between 0.33 and 0.66 the song is most likely to be a rap song. We can see this score as a way to quantify how melodic a song is. Therefore, we're going to keep this column to use as a predicting variable

```
[10]: spot.drop(index=spot[spot.speechiness > 0.66].index, inplace=True)
spot.shape
```

```
[10]: (138567, 19)
```

Now we're going to look into 'tempo' attribute We can notice that there are tracks whose tempo is zero, let's list them out

```
[11]: spot[spot.tempo == 0]
```

```
[11]:
```

	acousticness	artists	danceability	\
2721	0.099500	['Frank Sinatra']	0.0	
3387	0.756000	['Waylon Jennings']	0.0	
6930	0.931000	['Crain & Taylor']	0.0	
7411	0.111000	['Sound Dreamer']	0.0	
7792	0.145000	['Fan Sounds']	0.0	
...	
161443	0.916000	['Water Sound Natural White Noise']	0.0	
164201	0.862000	['Bill Cosby']	0.0	
169522	0.913000	['Granular']	0.0	
169743	0.000013	['Naturaleza FX']	0.0	
169770	0.957000	['Granular']	0.0	

	duration_ms	energy	explicit	id	\
2721	60280	0.906000	0	OP7TUyrm60fIDJJKcidvnu	
3387	14708	0.048400	0	2mex2o4uA69pMcLjMtyyGb	
6930	598425	0.000075	0	3oKBZhpwrMi0hosXauv3lP	
7411	5403500	0.000099	0	7foc25ig7dibxvULPU2kBG	
7792	500167	0.000020	0	4xu38KnBRHbRHRwdg4Kful	
...	
161443	63000	0.032000	0	5pGBDKBaR63vuJ4g8ialcU	
164201	215280	0.770000	0	2WOKFIFBFcLB1klD7ugiw6	
169522	205161	0.000164	0	2e6fCxt07NzsnujvliBtEk	
169743	150879	0.000020	0	4UFlnhDTGyKvlh0QziDHkG	
169770	146061	0.148000	0	4rkUd5Juj2icJoNTLq0jnP	

	instrumentalness	key	liveness	loudness	mode	\
2721	0.000018	1	0.366	-6.227	1	
3387	0.000144	4	0.166	-18.198	1	
6930	0.892000	1	0.115	-19.703	0	
7411	0.392000	2	0.137	-21.669	1	
7792	0.213000	6	0.114	-25.556	1	
...	
161443	0.202000	1	0.103	-30.704	1	
164201	0.000002	9	0.694	-15.316	0	
169522	0.910000	10	0.155	-31.221	0	
169743	0.208000	1	0.311	-16.873	1	
169770	0.168000	5	0.112	-22.012	1	

name popularity \

2721	My Kind Of Town (Reprise) - Live At The Sands ...	22
3387	Ride Me Down Easy	29
6930	Ocean Waves	47
7411	Brown Noise - 90 Minutes	50
7792	Box Fan Long Loop For Sleep	60
...
161443	Deep Sleep Recovery Noise	70
164201	Noah: Right!	16
169522	White Noise - 700 hz	64
169743	Colors of the Rain	68
169770	Brown Noise 750 LPF	69

	release_date	speechiness	tempo	valence	year
2721	1966-07	0.0	0.0	0.0	1966
3387	1973-07-01	0.0	0.0	0.0	1973
6930	2008-10-01	0.0	0.0	0.0	2008
7411	2013-06-05	0.0	0.0	0.0	2013
7792	2017-01-01	0.0	0.0	0.0	2017
...
161443	2020-02-25	0.0	0.0	0.0	2020
164201	1963	0.0	0.0	0.0	1963
169522	2017-10-14	0.0	0.0	0.0	2017
169743	2019-10-10	0.0	0.0	0.0	2019
169770	2019-01-11	0.0	0.0	0.0	2019

[95 rows x 19 columns]

Interestingly, by actually search for the track on Spotify and listen to it, we recognize that there are 2 groups of tracks whose tempo is zero. The first group are tracks of noises or sounds such as 'white noise' or 'water sound', which people listen to to focus or relax. The second group comprises tracks which are actually songs for example with tempo but for some reasons, Spotify engine was unable to detect. We assume these data entries are corrupted and going to remove from the dataset.

```
[12]: spot.drop(index=spot[spot.tempo == 0].index, inplace=True)
spot.shape
```

[12]: (138472, 19)

```
[13]: # Rename 'duration_ms' column to 'duration'
spot.rename(columns={'duration_ms': 'duration'}, inplace=True)
# We're going to drop columns not related to the analysis such as 'artists',
↳ 'id', 'release_date' & 'name', ...
spot.drop(columns=['artists', 'id', 'release_date', 'name', 'liveness',
↳ 'year'], inplace=True)
```

```
[14]: # Normalize 'loudness' column
from sklearn.preprocessing import MinMaxScaler
```

```
mms = MinMaxScaler()
spot.loc[:, 'loudness'] = mms.fit_transform(spot['loudness'].to_frame())
spot.describe()
```

```
[14]:
```

	acousticness	danceability	duration	energy	\
count	138472.000000	138472.000000	1.384720e+05	138472.000000	
mean	0.421318	0.548262	2.329282e+05	0.529960	
std	0.355001	0.173255	1.061279e+05	0.260163	
min	0.000000	0.055100	2.462700e+04	0.000000	
25%	0.062600	0.432000	1.750192e+05	0.323000	
50%	0.355000	0.557000	2.154400e+05	0.539000	
75%	0.773000	0.673000	2.661070e+05	0.743000	
max	0.996000	0.988000	4.270034e+06	1.000000	

	explicit	instrumentalness	key	loudness	\
count	138472.000000	138472.000000	138472.000000	138472.000000	
mean	0.087122	0.127792	5.200662	0.739011	
std	0.282015	0.279001	3.517606	0.097110	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	2.000000	0.688191	
50%	0.000000	0.000088	5.000000	0.754971	
75%	0.000000	0.025000	8.000000	0.810783	
max	1.000000	1.000000	11.000000	1.000000	

	mode	popularity	speechiness	tempo	\
count	138472.000000	138472.000000	138472.000000	138472.000000	
mean	0.710042	37.827922	0.071856	118.749887	
std	0.453744	18.095701	0.077200	30.184993	
min	0.000000	1.000000	0.000000	30.946000	
25%	0.000000	25.000000	0.033900	95.390750	
50%	1.000000	39.000000	0.043000	116.203500	
75%	1.000000	51.000000	0.068100	137.478000	
max	1.000000	100.000000	0.660000	244.091000	

	valence
count	138472.000000
mean	0.537731
std	0.261161
min	0.000000
25%	0.328000
50%	0.548000
75%	0.756000
max	1.000000

```
[15]: spot.shape
```


```
[15]: (138472, 13)
```

After cleaning the dataset we've arrived with a dataset having 12 features and 1 label column

```
[16]: # Create dependent and independent variables
X = spot[['acousticness', 'danceability', 'energy', 'explicit',
          ↪ 'instrumentalness', 'loudness', 'speechiness']].to_numpy()
y = spot.popularity.to_numpy()
```

```
[17]: spot.corr().popularity.plot(kind='bar', ylabel='Correlation')
```

```
[17]: <AxesSubplot:ylabel='Correlation'>
```




We can see there is a relatively strong correlation in 'acousticness', 'energy', 'loudness' and 'explicit'. Interestingly, 'valence' has close to zero bearing on a track's 'popularity', which means whether the theme of the song is 'happy' or 'sad' has little impact on its popularity. A track's 'duration', 'key' and 'mode' and 'tempo' also have close to no influence over popularity

Let's explore the relationship with popularity of each individual attributes of interest


```
[18]: s_lr = LinearRegression()
s_lr.fit(X[:, 0].reshape(-1, 1), y)
corr = spot[['acousticness', 'popularity']].corr().to_numpy()

ax = spot.plot(kind='scatter', x='acousticness', y='popularity', figsize=[10, 10])
x = np.arange(0, 1, 0.01)
y_ = s_lr.intercept_ + s_lr.coef_*x
ax.plot(x, y_, color='red')
ax.set_title('Correlation between \'acousticness\' and \'popularity\'')
ax.text(0.6, 50, 'Pearson correlation coeff = {:.2f}'.format(corr[0, 1]))
plt.show()
```




```
[19]: s_lr = LinearRegression()
s_lr.fit(X[:, 1].reshape(-1, 1), y)
corr = spot[['danceability', 'popularity']].corr().to_numpy()

ax = spot.plot(kind='scatter', x='danceability', y='popularity', figsize=[10, 10])
x = np.arange(0, 1, 0.01)
y_ = s_lr.intercept_ + s_lr.coef_*x
ax.plot(x, y_, color='red')
ax.set_title('Correlation between \'danceability\' and \'popularity\'')
ax.text(0.6, 50, 'Pearson correlation coeff = {:.2f}'.format(corr[0, 1]))
plt.show()
```



```
[20]: s_lr = LinearRegression()
s_lr.fit(X[:, 2].reshape(-1, 1), y)
corr = spot[['energy', 'popularity']].corr().to_numpy()


ax = spot.plot(kind='scatter', x='energy', y='popularity', figsize=[10, 10])
x = np.arange(0, 1, 0.001)
y_ = s_lr.intercept_ + s_lr.coef_*x
ax.plot(x, y_, color='red')
ax.set_title('Correlation between \'energy\' and \'popularity\'')
ax.text(0.2, 40, 'Pearson correlation coeff = {:.2f}'.format(corr[0, 1]))
plt.show()
```




```
[21]: s_lr = LinearRegression()
s_lr.fit(X[:, 3].reshape(-1, 1), y)
corr = spot[['explicit', 'popularity']].corr().to_numpy()


ex0 = spot.groupby('explicit').get_group(0)['popularity']
ex1 = spot.groupby('explicit').get_group(1)['popularity']

fig, ax = plt.subplots()
fig.set_size_inches(10, 10)
ax.boxplot(x=[ex0, ex1], labels=[0, 1])
ax.set_title('Correlation between \'loudness\' and \'popularity\'')
ax.text(1.2, 50, 'Pearson correlation coeff = {:.2f}'.format(corr[0, 1]))
plt.show()
```




```
[22]: s_lr = LinearRegression()
s_lr.fit(X[:, 4].reshape(-1, 1), y)
corr = spot[['instrumentalness', 'popularity']].corr().to_numpy()

ax = spot.plot(kind='scatter', x='instrumentalness', y='popularity',
    figsize=[10, 10])
x = np.arange(0, 1, 0.001)
y_ = s_lr.intercept_ + s_lr.coef_*x
ax.plot(x, y_, color='red')
ax.set_title('Correlation between \'instrumentalness\' and \'popularity\'')
ax.text(0.6, 50, 'Pearson correlation coeff = {:.2f}'.format(corr[0, 1]))
plt.show()
```




```
[23]: s_lr = LinearRegression()
s_lr.fit(X[:, 5].reshape(-1, 1), y)
corr = spot[['loudness', 'popularity']].corr().to_numpy()

ax = spot.plot(kind='scatter', x='loudness', y='popularity', figsize=[10, 10])
x = np.arange(0, 1, 0.001)
y_ = s_lr.intercept_ + s_lr.coef_*x
ax.plot(x, y_, color='red')
ax.set_title('Correlation between \'loudness\' and \'popularity\'')
ax.text(0.6, 50, 'Pearson correlation coeff = {:.2f}'.format(corr[0, 1]))
plt.show()
```



```
[24]: s_lr = LinearRegression()
s_lr.fit(X[:, 6].reshape(-1, 1), y)
corr = spot[['speechiness', 'popularity']].corr().to_numpy()

ax = spot.plot(kind='scatter', x='speechiness', y='popularity', figsize=[10, 10])
x = np.arange(0, 1, 0.001)
y_ = s_lr.intercept_ + s_lr.coef_*x
ax.plot(x, y_, color='red')
ax.set_title('Correlation between \'speechiness\' and \'popularity\'')
ax.text(0.6, 50, 'Pearson correlation coeff = {:.2f}'.format(corr[0, 1]))
plt.show()
```



We can see from the graphs above that the degree of correlation is not very apparent and a great

amount of noise is observed for each case. Now we're going to create train and test set using `train_test_split()` function from sklearn library

```
[25]: from sklearn.model_selection import train_test_split

# Split into train and test set
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
→random_state=1)

X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
[25]: ((96930, 7), (41542, 7), (96930,), (41542,))
```

Let's now first try to build a prediction model based on a Multiple Linear Regression method and evaluate its accuracy using test set, metrics for accuracy used here are Mean Absolute Error (MAE) score and R2 score

```
[26]: lr = LinearRegression()
lr.fit(X_train, y_train)
y_hat = lr.predict(X_test)

mean_absolute_error(y_test, y_hat), r2_score(y_test, y_hat)
```

```
[26]: (12.399115174516796, 0.28638612575513445)
```

Relatively decent score, on average, a prediction for the popularity score of a song is off by +- 12 points. However, we can try to improve evaluation score by trying a Polynomial Regression model

Grid search combined with cross validation method is going to be used to pick out the polynomial degree which yields the highest accuracy score.

```
[27]: pf = PolynomialFeatures(degree=2)
lr = LinearRegression()
pipe = Pipeline(steps=[('poly', pf), ('reg', lr)])

grid = {'poly__degree': [i for i in range(2, 6, 1)]}

pf.get_params()
search = GridSearchCV(pipe, grid)
search.fit(X_train, y_train)
print(search.best_score_, search.best_params_)
```

```
0.38643776545205477 {'poly__degree': 4}
```

A grid search with cross validation scoring return the best polynomial degree for the Polynomial Regression is 4. We're going to set the model parameter to degree=4 and test using test dataset

Results

```
[28]: pf = PolynomialFeatures(degree=4)
lr = LinearRegression()
pipe = Pipeline(steps=[('poly', pf), ('reg', lr)])

pipe.fit(X_train, y_train)
y_hat = pipe.predict(X_test)

print('Mean absolute score is: {:.2f}'.format(mean_absolute_error(y_test,
→y_hat)))
print('R2 score is: {:.2f}'.format(r2_score(y_test, y_hat)))
```

Mean absolute score is: 11.37

R2 score is: 0.38

Fairly consistent with scores obtained from cross validation. Now let's assume we're going to compose an EDM track which is, by Spotify model's definition, very danceable, high in energy and relatively loud with a bit of rap. Let's see how well this song can do

```
[29]: # X = spot[['acousticness', 'danceability', 'energy', 'explicit',
→'instrumentalness', 'loudness', 'speechiness']].to_numpy()

test = [[0.1, 0.6, 0.8, 1, 0.5, 0.9, 0.2]]

print(pipe.predict(test))
```

[52.10633809]

Such a song is predicted to receive a score of 51, which is higher than the mean popularity score of the dataset. Not bad!

Discussion

An R2 score of around 0.38 indicates that about 38% of the variance of the values in 'popularity' column is explained by the model. Given that this model involves human behaviours and it is in general difficult to explain why or why not a person like something, a 0.38 R2 score is acceptable as to the model's power to predict.

It is important to note that this model does not take into account the effect of the popularity of the artist, the size of such artist fanbase (and how much of it are Spotify users), or how a track is marketed. These elements could be a deciding factor as to whether a track can gain the initial traction to become more popular.

Another point to note is that this model accuracy depends greatly on how well or how accurately the Spotify machine learning model quantifies attributes of a track. For instance, if Spotify model gives a track a 'energy' score of 1 but, in reality, the general perception of the track by the majority of Spotify users who have listened to the song is that it is not very energetic then this data entry will introduce noise into the model.

In order to improve this prediction model, more data about how tracks are recommended to Spotify users or how many times a track is listed under different playlists might prove to be useful.

Conclusion

In this report we have presented an analysis on the Spotify tracks dataset. We found that there's noticeable correlation between a track's 'popularity' and its attributes such as 'loudness', 'energy' and 'acousticness'. A predictive model using Polynomial Regression was then built to predict a song's popularity on Spotify based on its intrinsic characteristics. The model achieved a reasonable R^2 score of 0.38.